

# Practical Deep RL

**Volodymyr Mnih, DeepMind**

UCL

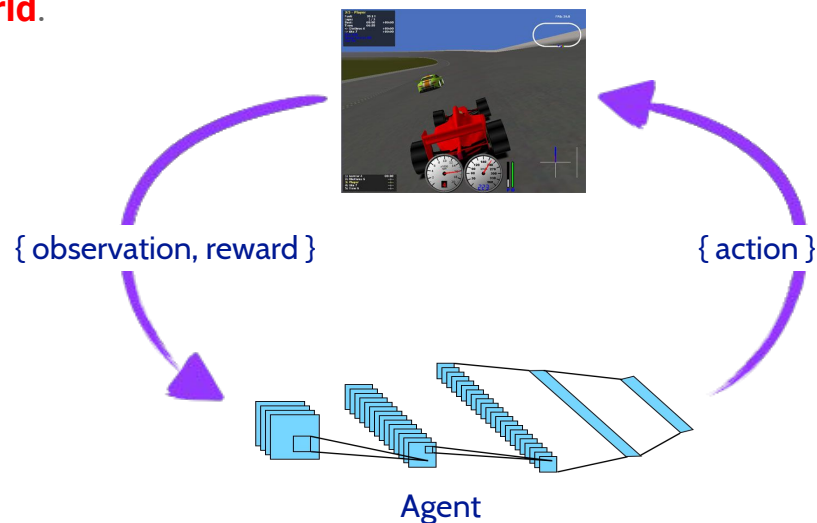
March 16, 2017



DeepMind


# Deep Reinforcement Learning == AI?

- AI - building machines that are good at **sequential decision-making problems humans care about**.
- **Reinforcement learning** is a general framework for studying sequential decision-making.
- Deep learning:
  - Current best way of making computers **perceive the world**.
  - More generally it is a framework for learning in deep computational graphs.



# The Deep RL Boom

TD-Gammon (Tesauro, 1989-1995)




Slot car driving (Lange & Riedmiller, 2012)




Arcade Learning Environment (Bellemare et al, 2013)




Deep Q-Networks (2013, 2015)



Trust region policy optimization (Schulman et al, 2015)



End-to-end training on real robots (Levine et al, 2015)



# Deep RL vs Deep Learning


- How is deep RL different from standard deep learning?
- The data distribution is non-stationary.
  - Neural nets don't like this. Most of the theory no longer applies.
- The data distribution is determined by the agent's actions.
  - Exploration vs exploitation.
  - You can get stuck in local minima. Optimization really matters.
- Sparse/delayed feedback.
- Training neural nets with RL was thought to be inherently unstable (Tsitsiklis & Van Roy, 1997).

# Outline

- DQN in more detail.
- Faster agents through parallel training.
- Better data efficiency through unsupervised RL.
- Some practical advice.

# Deep Q Networks (DQN)

- Represent the action value (Q) function using a convolutional neural network.
- Train using end-to-end Q-learning.
- Can we do this in a stable way?



# DQN

Initialize **target network**  $\theta^- \leftarrow \theta$

For each time step  $t$

Take action  $a_t$ , and observe  $r_t, s_{t+1}$

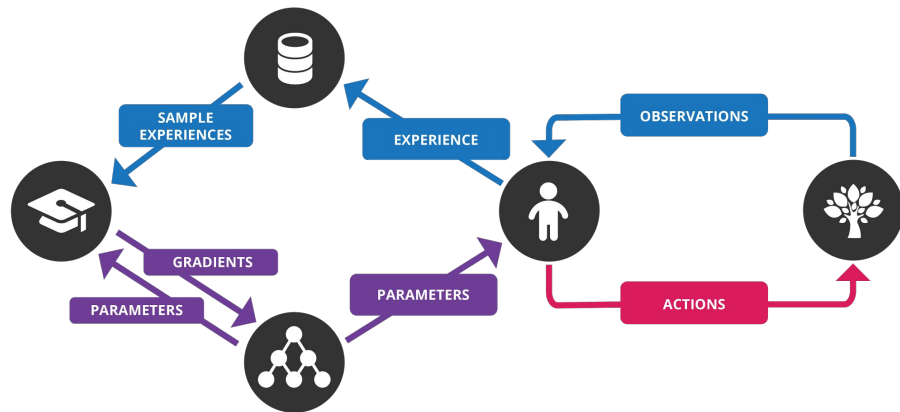
**Sample**  $(s, a, r, s')$  from **replay memory**

Generate **target**  $r + \delta \gamma \max_{a'} Q(s', a'; \theta^-)$

Take SGD step following  $\theta_{t+1} \leftarrow \theta_t - \eta \frac{\partial L(\theta)}{\partial \theta_t}$

Update **target network** if  $t \% k : \theta^- \leftarrow \theta$

**Store**  $(s_t, a_t, r_t, s_{t+1})$  in **replay memory**



# DQN

- High-level idea - make Q-learning look like supervised learning.
- Apply Q-updates on batches of past experience instead of online:
  - Experience replay (Lin, 1993).
  - Previously used for better data efficiency.
  - Makes the data distribution more stationary.
- Use an older set of weights to compute the targets (**target network**):
  - Keeps the target function from changing too quickly.


$$L_i(\theta_i) = \mathbb{E}_{s,a,s',r \sim D} \left( \underbrace{r + \gamma \max_{a'} Q(s', a'; \theta_i^-)}_{\text{target}} - Q(s, a; \theta_i) \right)^2$$



# Target Network Intuition

- Changing the value of one action will change the value of other actions and similar states.
- The network can end up chasing its own tail because of bootstrapping.
- Somewhat surprising fact - bigger networks are less prone to this because they alias less.

$$L_i(\theta_i) = \mathbb{E}_{s,a,s',r \sim D} \left( \underbrace{r + \gamma \max_{a'} Q(s', a'; \theta_i^-)}_{\text{target}} - Q(s, a; \theta_i) \right)^2$$



# Neural Fitted Q Iteration

- NFQ (Riedmiller, 2005) trains neural networks with Q-learning.
- Alternates between collecting new data and fitting a new Q-function to all previous experience with batch gradient descent.


```
NFQ_main() {  
  input: a set of transition samples  $D$ ; output: Q-value function  $Q_N$   
  k=0  
  init_MLP()  $\rightarrow Q_0$ ;  
  Do {  
    generate_pattern_set  $P = \{(input^l, target^l), l = 1, \dots, \#D\}$  where:  
       $input^l = s^l, u^l$ ,  
       $target^l = c(s^l, u^l, s'^l) + \gamma \min_b Q_k(s'^l, b)$   
    Rprop_training( $P$ )  $\rightarrow Q_{k+1}$   
    k:= k+1  
  } WHILE ( $k < N$ )
```

- DQN can be seen as an online variant of NFQ.


# Lin's Networks

- Long-Ji Lin's thesis "Reinforcement Learning for Robots using Neural Networks" (1993) also trained neural nets with Q-learning.
- Introduced experience replay among other things.
- Lin's networks did not share parameters among actions.

**Lin's architecture**




**DQN**



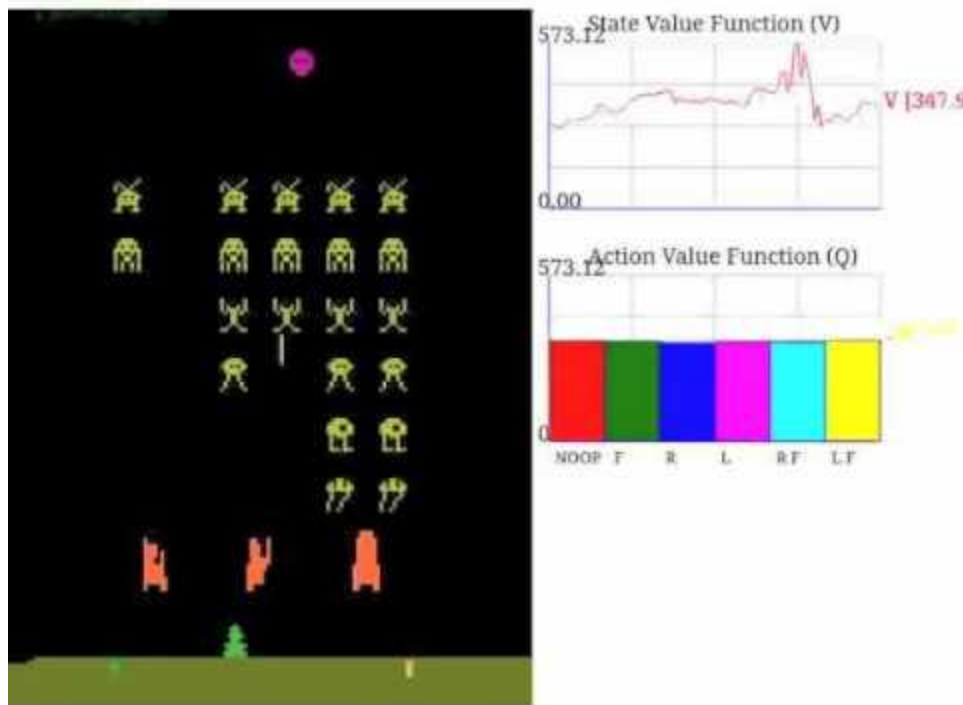
# DQN Playing ATARI



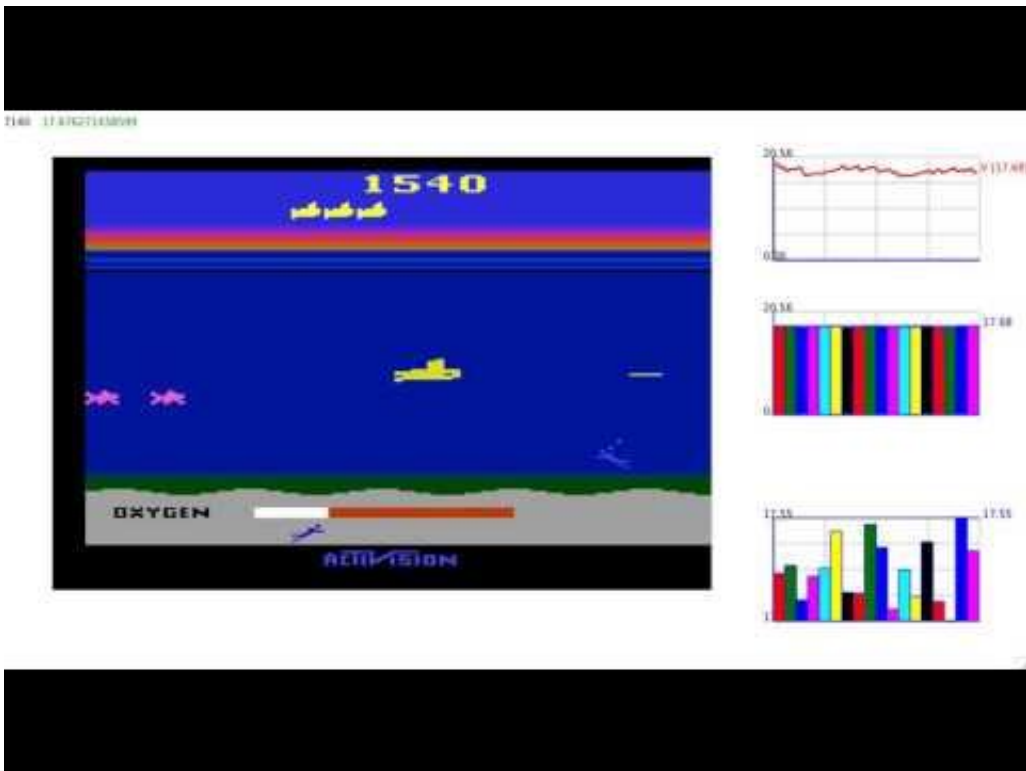
# Action Values on Pong



# Learned Value Functions




# Sacrificing Immediate Rewards




# DQN Atari Filters

- Visualizing the convolutional filters learned by DQN.
- Surprisingly little structure compared to convolutional filters on ImageNet.




ImageNet filters, Krizhevsky et al. (2012)





# Labyrinth Filters

- Filters learned by A3C on a 3D environment.
- Visually richer environment produces more structured and interesting filters.




Labyrinth agent filters



ImageNet filters, Krizhevsky et al. (2012)

# t-SNE



- Problem - visualizing high dimensional data by projecting it to 2D.
- Multi-Dimensional Scaling (MDS):
  - Project to 2D while maintaining pairwise distances.
  - Crowding problem - impossible to maintain all distances.
- t-distributed Stochastic Neighbor Embedding (Van Der Maaten and Hinton, 2008).
  - Try to preserve neighbors but allow some nearby points in the high dimensional space to be far apart in the low dimensional space.
  - Nearby points in a t-SNE plot are nearby in the original space.
  - Most distant points in a t-SNE plot are distant in the original space, but a few could be nearby.



Gashler et al., 2011


# Space Invaders Representation

- t-SNE embedding of the last hidden layer of the DQN network.


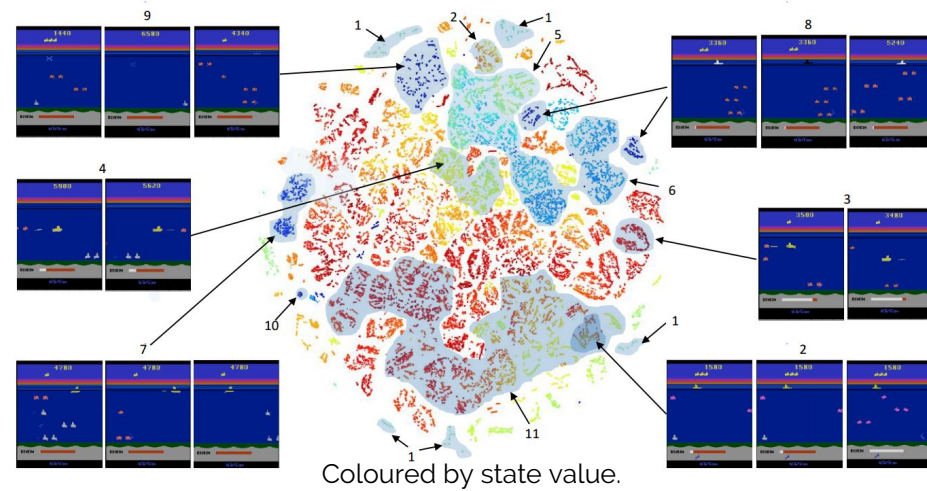


# Understanding DQN (Zahavy et al., 2015)

- More detailed t-SNE analysis of DQN.
- Argues that DQN learns hierarchical state aggregation and hierarchical policies.





Coloured by remaining oxygen.



# Saliency Maps

- Value-Advantage decomposition of Q:  
$$Q^{\pi}(s, a) = V^{\pi}(s) + A^{\pi}(s, a)$$
- Dueling DQN (Wang et al., 2015):




# Beyond DQN

- DQN is very robust, but has some limitations:
  - Computationally expensive.
    - But see Prioritized DQN (Schaul et al., 2016), Dueling Prioritized DQN (Wang et al, 2016)
  - Has not worked as well with recurrent networks.
  - Does not handle continuous actions (see DDPG Lillicrap et al., 2015).
- Is there a method that allows:
  - Fast training on a single machine - hours for simple Atari games.
  - Freedom to use on or off-policy methods.
  - Flexibility - discrete or continuous actions, feedforward or recurrent models, etc.

# AsyncRL

- Asynchronous training of RL agents:
  - Parallel actor-learners implemented using **CPU threads** and shared parameters.
  - Online **Hogwild!**-style asynchronous updates (Recht et al., 2011, Lian et al., 2015).
  - No replay? Parallel actor-learners have a similar stabilizing effect.
  - Choice of RL algorithm: on-policy or off-policy, value-based or policy-based.





# 1-step Q-Learning

- Parallel actor-learners compute online 1-step update


$$y \leftarrow r + \gamma \max_{a'} Q(s', a'; \theta^-)$$
$$\Delta\theta \leftarrow \Delta\theta + \frac{\partial (y - Q(s, a; \theta))^2}{\partial \theta}$$

- Gradients accumulated over minibatch before update



# N-step Q-Learning

- Q-learning with a uniform mixture of backups of length 1 through N.




$$y \leftarrow \sum_{k=0}^{N-1} \gamma^k r_{t+k} + \gamma^N \max_{a'} Q(s_{t+N}, a'; \theta^-)$$

$$\Delta \theta \leftarrow \Delta \theta + \frac{\partial (y - Q(s_t, a_t; \theta))^2}{\partial \theta}$$

- Variation of “Incremental multi-step Q-learning” (Peng & Williams, 1995).

# Async Advantage Actor-Critic (A3C)

- The agent learns a policy and a state value function.
- Policy gradient multiplied by an estimate of the advantage. Similar to Generalized Advantage Estimation (Schulman et al, 2015).



$$\nabla_{\theta} \log \pi(a_t | s_t, \theta) \left( \sum_{k=0}^N \gamma^k r_{t+k} + \gamma^{N+1} V(s_{t+N+1}) - V(s_t) \right)$$

# Async Advantage Actor-Critic (A3C)

---

**Algorithm S3** Asynchronous advantage actor-critic - pseudocode for each actor-learner thread.

---

*// Assume global shared parameter vectors  $\theta$  and  $\theta_v$  and global shared counter  $T = 0$*

*// Assume thread-specific parameter vectors  $\theta'$  and  $\theta'_v$*

Initialize thread step counter  $t \leftarrow 1$

**repeat**

Reset gradients:  $d\theta \leftarrow 0$  and  $d\theta_v \leftarrow 0$ .

Synchronize thread-specific parameters  $\theta' = \theta$  and  $\theta'_v = \theta_v$

$t_{start} = t$

Get state  $s_t$

**repeat**

Perform  $a_t$  according to policy  $\pi(a_t|s_t; \theta')$

Receive reward  $r_t$  and new state  $s_{t+1}$

$t \leftarrow t + 1$

$T \leftarrow T + 1$

**until** terminal  $s_t$  **or**  $t - t_{start} == t_{max}$

$R = \begin{cases} 0 & \text{for terminal } s_t \\ V(s_t, \theta'_v) & \text{for non-terminal } s_t // \text{ Bootstrap from last state} \end{cases}$

**for**  $i \in \{t - 1, \dots, t_{start}\}$  **do**

$R \leftarrow r_i + \gamma R$

Accumulate gradients wrt  $\theta'$ :  $d\theta \leftarrow d\theta + \nabla_{\theta'} \log \pi(a_i|s_i; \theta')(R - V(s_i; \theta'_v))$

Accumulate gradients wrt  $\theta'_v$ :  $d\theta_v \leftarrow d\theta_v + \partial (R - V(s_i; \theta'_v))^2 / \partial \theta'_v$

**end for**


Perform asynchronous update of  $\theta$  using  $d\theta$  and of  $\theta_v$  using  $d\theta_v$ .

**until**  $T > T_{max}$

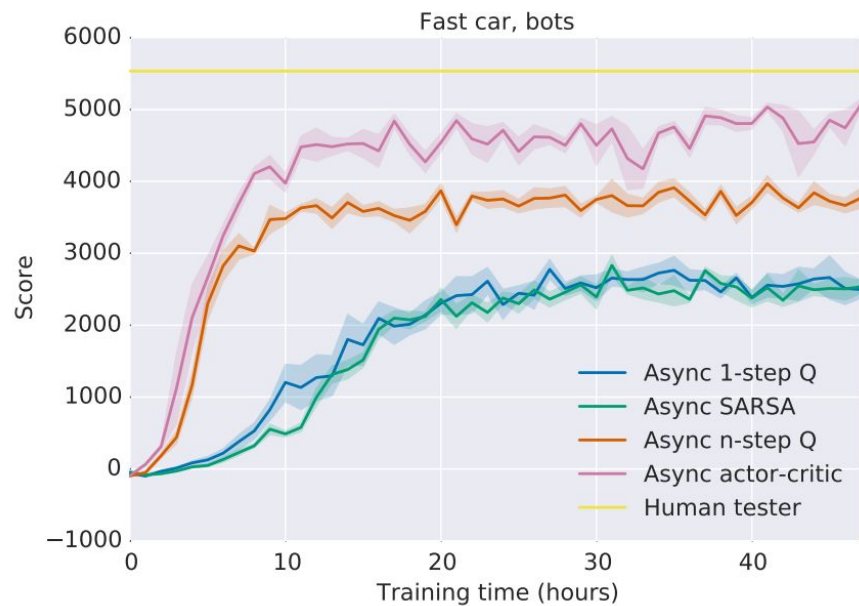

---

# AsyncRL - Learning Speed

- New asynchronous methods trained on 16 CPU cores compared to DQN (blue) trained on a K40 GPU.
- n-step methods can be much faster than single step methods.
- Async advantage actor-critic tends to dominate the value-based methods.




# AsyncRL - TORCS




# AsyncRL - Scalability

- Average speedup from using K threads to reach a reference score averaged over 7 Atari games.
- **Super-linear** speed-up for 1-step methods.




# Data Efficiency of 1-Step Q-learning

- Better **data efficiency** from multiple actor-learners plus a speedup from parallel training.
  - 1 thread (blue) 16 threads (yellow)



# Data Efficiency of A3C

- No data-efficiency gains. Sub-linear speedup from parallel training.
  - 1 thread (blue) 16 threads (yellow)







# A3C - ATARI Results

Method	Training Time	Mean	Median
DQN	8 days on GPU	121.9%	47.5%
Gorilla	4 days, 100 machines	215.2%	71.3%
D-DQN	8 days on GPU	332.9%	110.9%
Dueling D-DQN	8 days on GPU	343.8%	117.1%
Prioritized DQN	8 days on GPU	463.6%	127.6%
<b>A3C, FF</b>	1 day on CPU	344.1%	68.2%
<b>A3C, FF</b>	4 days on CPU	496.8%	116.6%
<b>A3C, LSTM</b>	4 days on CPU	623.0%	112.6%

# A3C and Data Efficiency

- Speed may not be the best metric.
- A3C is not very data efficient.
- ACER (Wang et al., 2017) and PGQ (O'Donoghue et al., 2017) combine replay with A3C for improved data efficiency.





# A3C - Procedural Maze Navigation in 3D



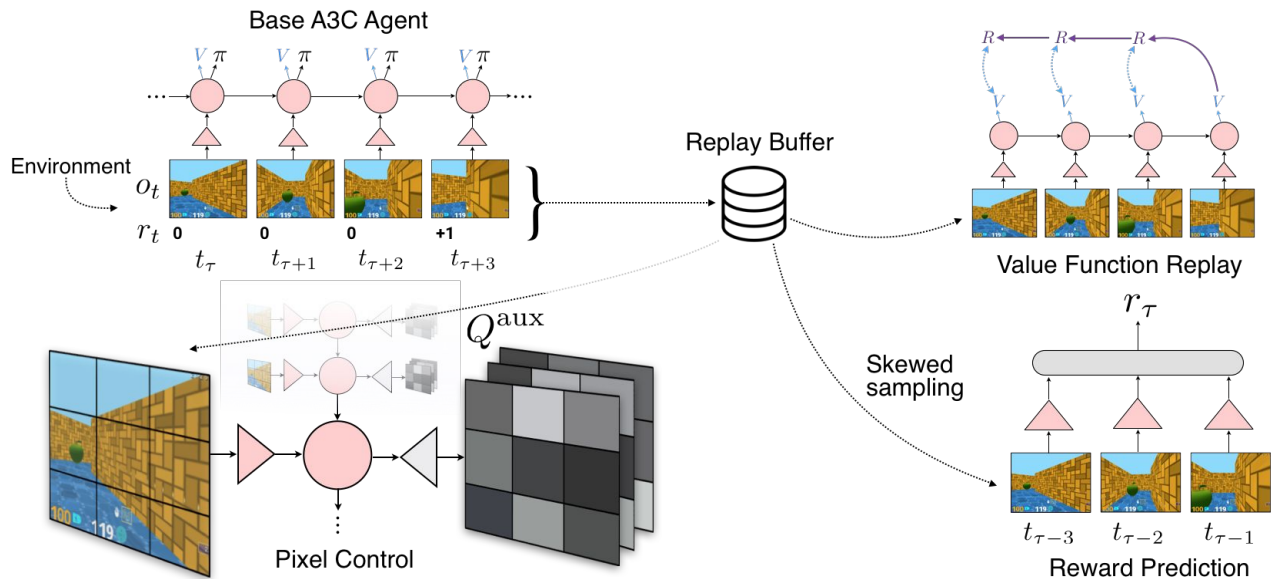
# Unsupervised Reinforcement Learning

- The best deep RL methods are still very data hungry. Especially with **sparse rewards**.
- Obvious solution - Learn about the environment.
- We can augment an RL agent with **auxiliary prediction and control tasks** to improve data efficiency.
- The UNREAL agent - UNsupervised REinforcement and Auxiliary Learning.
  - “Reinforcement Learning with Unsupervised Auxiliary Tasks”, Jaderberg et al. (2017)




# The UNREAL Architecture

- UNREAL augments an LSTM A3C agent with 3 auxiliary tasks.
- Can be used on top of DQN, DDPG, TRPO or other agents.



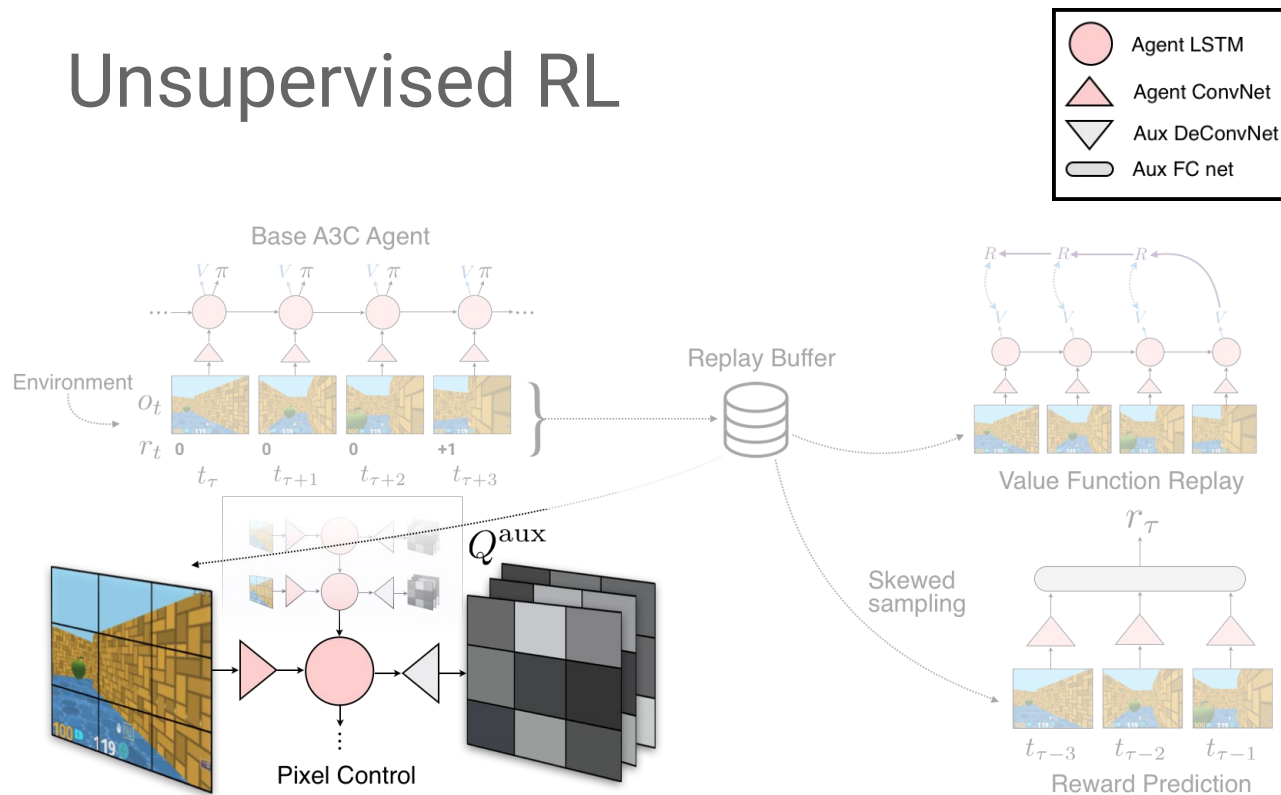
# The UNREAL Architecture

- Base A3C LSTM agent learns from the environment's scalar reward signal.
- UNREAL acts using the base A3C agent's policy.



# Unsupervised RL


- Augment A3C with many **auxiliary control tasks**.
- Learning to control many aspects of the environment.
- Pixel control - learn to maximally change parts of the screen.
- Feature control (not used by UNREAL) - learn to control the internal representations.



# The UNREAL Architecture

Focusing on rewards:

- Rebalanced reward prediction.
- Shape the agent's CNN by classifying whether a sequence of frames will lead to reward.
- No need to worry about off-policy learning.

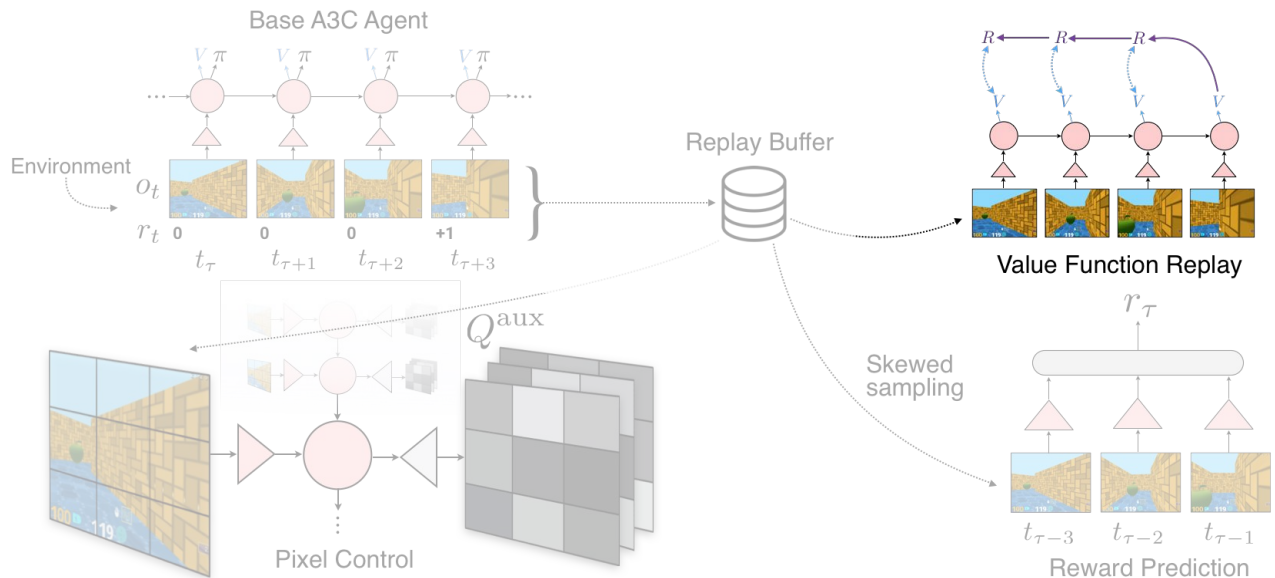




# The UNREAL Architecture

Focusing on rewards:


- Value function replay.
- Faster learning of the value function.






# DeepMind Lab Results



- Average human-normalized performance on 13 3D environments from DeepMind Lab.
- Tasks include random maze navigation and laser tag.
- Roughly a 10x improvement in data efficiency over A3C.
- 60% improvement in final performance.




# Feature Control



# Unsupervised RL Baselines



# Montezuma's Revenge



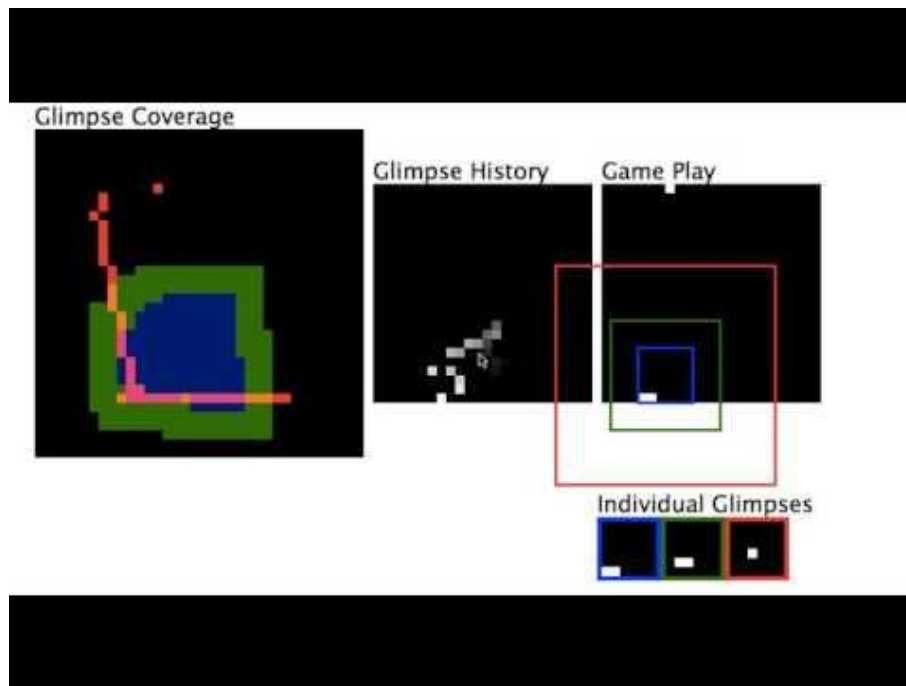
# UNREAL playing DeepMind Lab



# Practical Advice - Getting Started

- Start with a simple problem.
  - Something solvable in under a minute on your local machine.
  - Make it similar to the problem you really want to solve.
  - Ideally it should have knobs for controlling its difficulty.
- Plot the training curves (averaged over multiple episodes).
- Visualize the policy.
- Visualize the value function.
- Visualize everything you can think of.

# The game of Catch





# Practical Advice - Neural Nets

- Doing early experiments with a small network can help iterate faster.
  - This can also backfire (DQN and target networks).
- Reasonable strategy:
  - Run a few progressively larger nets to find what's sufficient for experimenting.
  - Periodically try larger nets to max out performance and verify assumptions.
- Be careful with initialization:
  - Visualize the initial policy to make sure it gets some rewards.
- Try RMSProp and/or Adam.
- Test deep learning tricks before incorporating them: dropout, batch norm, etc.
- See John Schulman's excellent guide - <http://joschu.net/docs/nuts-and-bolts.pdf>

# Questions?